


Parallel Programming | ParProg

Zusammenfassung

1. MULTI-THREADING

Parallelism (Subprograms run simultaneously for faster programs) vs. **Concurrency** (interleaved execution of programs for simpler programs)



Process: Program under execution, own address space (heavy-weight). **Proc:** Process execution and responsiveness, Cons: Interprocess communication overhead, expensive in creation, slow context switching and process termination.

Thread: Parallel sequence within a process. Sharing the same address space, but separate stack and registers (lightweight because most of the overhead already happened in the process creation).

Multi-threads: Changes made by one thread to shared resources will be seen by other threads.

Context switch: Required when changing threads. **Synchronous** (Thread waiting for condition) or **Asynchronous** (Thread gets released after defined time)

Multitasking: Cooperative (Threads must explicitly initiate context switches, scheduler can't interrupt) or Preemptive (scheduler can asynchronously interrupt thread via timer interrupt)

JVM Thread Model: JVM is a process in the OS. It runs as long as threads are running (Exception: threads marked as done with `setDaemon(true)` will not be waited upon). Threads are overlaid by the **Thread class** and the **Interface Runnable**. Code to be run in a Thread is within a overridden `run()`.

1.1. STARTING A THREAD

As an anonymous function (Lambda):

```
var myThread = new Thread(() -> { /* thread behaviour */ }); myThread.start();
```

As a named function:

```
var myThread = new Thread(() -> AsyncFunction(myThread)); myThread.start();
```

With explicit Runnable implementation:

```
class MyThread implements Runnable {
    @Override
    public void run() { /* thread behavior */ }
    var myThread = new Thread(new MyThread());
    myThread.start();
}
```

In C#: `var myThread = new Thread(() => { ... }); myThread.Start(); ... myThread.Join();`

Multi-Thread Example (no synchronization)

```
public static void main(String[] args) {
    var a = new Thread(() -> multiPrint("A"));
    var b = new Thread(() -> multiPrint("B"));
    System.out.println("Threads start"); a.start(); b.start(); // ...
    a.join(); b.join(); System.out.println("Threads joined");
}

static void multiPrint(String label) {
    for (int i = 0; i < 10; i++) { System.out.println(label + " " + i);
}
```

The printout of this function varies, it can be all possible combinations of A's and B's due to the non-deterministic scheduler

Thread Join: Waiting for a thread to finish (thread.join() blocks as long as thread is running).

```
a = new Thread(() -> multiPrint("A"));
b = new Thread(() -> multiPrint("B"));
System.out.println("Threads start"); a.start(); b.start(); // ...
a.join(); b.join(); System.out.println("Threads joined");
```

Thread States: Blocked (Thread is blocked and waiting for a monitor lock), New (Thread has not yet started), Runnable (Thread is runnable (Ready to run) normally), Terminated (Thread is terminated), Timed_Waiting

(Thread is waiting with a specified waiting time Thread.sleep(s)/Thread.join(s)), Waiting (Thread is waiting)

Yield: Thread is done processing for the moment and hints to the scheduler to release the processor. The scheduler can ignore this. Thread enters into ready-state. (Thread.yield())

Interrupts: Threads can also be interrupted from the outside (Thread.interrupt(), Thread can decide whether to resume or not) or from the inside (Thread.sleep() or Thread.join() can be interrupted). An `InterruptedException` is thrown. Otherwise a flag is set that can be checked with `interrupted()` or `isInterrupted()`

Exceptions: Exceptions thrown in `run()` can't be propagated to the Main thread. The exception needs to be handled within the code executed on the thread.

Thread Methods: `currentThread()`: (Reference to current thread, `setDaemon(true)`: (Mark a daemon thread), `getId()`/`getName()`: (Get thread ID/name), `isAlive()`: (Tests if thread is alive), `getStackTrace()`: (Get thread stack)

2. MONITOR SYNCHRONIZATION

Threads run arbitrarily. **Restriction of concurrency** for deterministic behavior.

Communication between threads: Sharing access to fields and the objects they refer to. Efficient, but poses problems: **Thread interference** and **memory consistency errors**.

Critical Section: Part of the code which must be executed by only 1 thread at a time for the values to stay consistent. Implementation with **mutual exclusion**.

synchronized: Body of method with the `synchronized` keyword is a critical section. Guarantees **memory consistency** and a **happens-before** relationship. Impossible for two invocations of a synchronized method on the same object to interleave. Other threads are **blocked** until the current thread is done with the object. Every object has a `Lock (Monitor/Lock)`. Maximum 1 thread can acquire the lock. Entry of a synchronized method acquires the lock of the object, it exits it.

public synchronized void deposit(int amount) { this.balance += amount; }

Can also be used within a method, the object that should be locked must be specified.

`synchronized(this) { this.balance += amount; }`

Exit synchronized block: End of the block, return, unhandled exceptions

2.0.1. Monitor Lock

A monitor is used for **internal mutual exclusion**. Only one thread operates at a time in Monitor. All non-private methods are synchronized. Threads can wait in **Monitor** for condition to be fulfilled. Can be inefficient with different waiting conditions, has **fairness-problems** and **no shared locks**.

Recursive Lock: A thread can acquire the same lock through recursive calls. Lock will be free by the last release.

Busy Wait: Running yield or sleep in a loop doesn't release the lock and is inefficient. Use wait/wait(). Waits on a condition. Temporarily releases Monitor-Lock so that other threads can run. Needs to be wrapped into a **while loop** to check if the wake up condition has been met.

WakeUp signal: Signalling a condition/thread in Monitor. `notify()` signals any waiting thread (sufficient if all threads wait for the same thing, it does not matter which ones are next - uniform waits or if only one single thread can continue like in a summing), `notifyAll()` wakes up all threads (e.g. one deposit can satisfy multiple withdraws, does not guarantee fairness). If a thread is woken up, it goes from the inner waiting room (waiting on a condition) into the **outer waiting room** (Thread has not started yet) where it waits for entry to the Monitor. There is no shortcut.

`IllegalMonitorStateException` is thrown if `notify`, `notifyAll` or `wait` is used outside synchronized.

```
// Java
class BankAccount {
    private int balance = 0;
    // Entry in the monitor
    public synchronized void withdraw
    (int amount) {
        while (amount > balance) { // not if
            wait(); // wait on a condition
        }
        balance -= amount;
        // release / leave monitor
        public synchronized void deposit
        (int amount) {
            balance += amount;
            notifyAll(); // Wakes up all waiting
        }
        threads in monitor inner waiting area
    }
}
```

3. SPECIFIC SYNCHRONIZATION PRIMITIVES

3.1. SEMAPHORE

Allocation of a limited number of free resources. Is in essence a **counter**. If a resource is **acquired**, count--, if a resource is **released**, count++. Can wait until resource becomes available. Can also acquire/release multiple permits at once atomically.

```
public class Semaphore {
    private int value; public Semaphore(int initial) { value = initial; }
    public synchronized void acquire() throws InterruptedException {
        while (value <= 0) { wait(); } value--;
    }
    public synchronized void release() { value++; notify(); }
}
```

General Semaphore: new Semaphore(`N`) (Counts from 0 to `N` for limited permits to access a resource)

Binary Semaphore: new Semaphore(`1`) (Counter for 1 if exclusive execution, open/closed)

Fair Semaphore: new Semaphore(`N`, `true`) (Uses FIFO Queue aging for fairness. Slower than non-fair variant)

Semaphores are **powerful**, any synchronization can be implemented. But relatively **low-level**.

```
class BoundedBuffer<T> {
    private Queue<T> queue = new LinkedList<>();
    private Semaphore upperLimit = new Semaphore(capacity, true); // how many free?
    private Semaphore lowerLimit = new Semaphore(0, true); // how many full?
    public void put(T item) throws InterruptedException {
        upperLimit.acquire(); // No. of free places - 1
        synchronized (queue) { queue.add(item); }
        lowerLimit.release(); // No. of full places + 1
    }
    public void get(T item) throws InterruptedException {
        T item;
        lowerLimit.acquire(); // No. of full places - 1
        synchronized (queue) { item = queue.remove(); }
        upperLimit.release(); // No. of free places + 1
        return item;
    }
}
```

3.2. LOCK & CONDITION

Monitor with multiple **waiting lists** and conditions. Independent from Monitor locks.

Lock-Object: Object for entry in the monitor (outer waiting room)

Condition-Object: Wait & Signal for a specific condition (inner waiting room).

ReentrantLock: Class in Java, **alternative to synchronized**. Allows multiple locking operations by the same thread and supported nested locking (Thread is able to re-enter the same lock).

Condition: Factors out the Object monitor methods `wait`, `notify` and `notifyAll` into distinct objects to give the effect of having multiple wait-sets per object, by combining them with the use of arbitrary Lock implementations. A **Condition** replaces the use of the **Object monitor methods**.

condition.await(): Throws an `InterruptedException` if the current thread has its interrupted status set on entry to this method or is interrupted while waiting (finally frees the lock in case of interrupt).

Buffer with Lock & Condition

```
class BoundedBuffer<T> {
    private Queue<T> queue = new LinkedList<>();
    private Lock monitor = new ReentrantLock(true); // fair queue
    private Condition nonFull = monitor.newCondition();
    private Condition nonEmpty = monitor.newCondition();
    ...
    public void put(T item) throws InterruptedException {
        monitor.lock(); // Lock queue
        try { while (queue.size() == capacity) { nonFull.await(); }
            queue.add(item); nonEmpty.signal(); } finally { monitor.unlock(); }
    }
    ...
    public T get() throws InterruptedException {
        monitor.lock(); // wait for queue to be filled & signal to other queue
        try { while (queue.size() == 0) { nonEmpty.await(); }
            T item = queue.remove(); nonFull.signal(); return item; }
        finally { monitor.unlock(); } // always release Lock, even after Exception
    }
}
```

3.3. READ-WRITE LOCK

Mutual exclusion is **unnecessary for read-only** threads. So one should allow parallel reading access, but implement mutual exclusion for write access.

```
ReadWriteLock readLock = new ReentrantReadWriteLock(true); // true for fairness
rLock.readLock().lock(); // shared lock
// read-only accesses here
rLock.readLock().unlock(); // release shared lock
// write (and read) accesses here
rLock.writeLock().lock(); // exclusive lock
rLock.writeLock().unlock(); // always exclusive lock
```

3.4. COUNT DOWN LATCH

Synchronization with a counter that can only **count down**. Threads can wait until counter ≤ 0 , or they can count down. The Latches can only be used once.

```
var ready = new CountDownLatch(N); var start = new CountDownLatch(1);
```

```
ready.countDown(); // wait for N cars ready.start(); // wait for all cars ready
```

```
start.countDown(); // wait for all cars ready.start(); // start the race
```

3.5. CYCLIC BARRIER

Meeting point for fixed number of threads. Threads wait until **everyone arrives**. Is **reusable**, threads can synchronize in multiple rounds at the same barrier (simplifies example above).

```
var start = new CyclicBarrier(N); start.await(); // all cars race as they're here
```

3.6. EXCHANGER

Rendez-Vous: Barrier with **information exchange** for 2 parties. Without exchange: **new CyclicBarrier(2)**, with exchange: **Exchanger.exchange(something)**. The Exchanger blocks until another thread also calls exchange(), returns argument x of the other thread.

```
var exchanger = new Exchanger<>();
for (int count = 0; count < 2; count++) { // odd number of exch.: last one blocks
    new Thread(() -> {
        for (int i = 0; i < 5; i++) {
            try {
                int out = exchanger.exchange(in);
                System.out.println(Thread.currentThread().getName() + " = got " + out);
            } catch (InterruptedException e) { }
        }
    }).start();
}
```

4. CONCURRENCY HAZARDS

4.1. RACE CONDITIONS

Insufficiently synchronized access to shared resources. The order of events affects the correctness of the program. Leads to **non-deterministic behavior**. Can occur without data race, but data race is often the cause.

Race Condition without data race: The critical section is not protected. Data Race is eliminated using synchronization, but there is no synchronization over larger blocks, so race conditions are still possible (e.g. non-atomic incrementing).

4.2. DATA RACE

Two threads in a single process access the **same variable** concurrently without synchronization, at least one of them is a write access.

Synchronize Everything? May not help and is expensive. So no.

4.3. THREAD SAFETY

Disposable classes in synchronization: Immutable Classes (Declaring all fields private and final and don't provide setters. Read-only Objects (Read-only access are thread-safe).

Confinement: Object belongs to only one thread at a time. **Thread Confinement** (Object belongs to only one thread), **Object Confinement** (Object is encapsulated in other synchronized objects)

Thread safe: A data type or method that behaves correctly when used from multiple threads as if it was running in a single thread without any additional coordination (Java concurrenc collections).

Thread Safety: Avoidance of Data Races. When no sharing is intended, give each thread a private copy of the data. When sharing is important, provide explicit synchronization.

4.4. DEADLOCKS

Happens when threads lock each other out, prohibiting both from running. Programs with potential deadlock are not considered correct. Threads can suddenly block each other.

```
// Thread 1
synchronized(ListA) {
    synchronized(ListB) {
        ListA.addAll(ListA);
    }
}
// Thread 2
synchronized(ListB) {
    synchronized(ListA) {
        ListA.addAll(ListB);
    }
}
```

Both threads in this scenario have **locked each other** out, the program cannot continue.

LiveLock: Threads have blocked each other permanently, but still execute wait instructions and therefore consume CPU during deadlock.

```
// Thread 1
b = false; while(a) { ... b = true;
// Thread 2
a = false; while(b) { ... a = true;
```

4.5. RESOURCE GRAPH

Thread T waits for Lock Thread T requires Lock of Resource R



Deadlocks can be identified by **cycles in the resource graph**.

Deadlock Avoidance: Introduce **linear blocking order**, lock nested only in ascending order. Or use **coarse granular locks** (Used when ordering does not make sense, e.g. block the whole bank instead of lock all accounts)

4.6. STARVATION

A thread never gets chance to access a resource. **Avoidance:** Use fair synchronization constructs. (Aging, stable fairness in previous synchronization constructs. Monitor and Thread priorities have a fairness problem.)

4.6. PARALLELISM CORRECTNESS CRITERIA

Safety: No race conditions and no deadlocks, **Liveness:** No starvation

4.7. .NET SYNCHRONIZATION PRIMITIVES

Monitor with sync object: `private object sync = new Object(); lock(sync) { ... }`

Use `Monitor.Acquire(sync)`, `Monitor.Release(sync)`. Uses fair FIFO-Queue. Locks: No fairness flag, no Lock & Condition. **Additional:** `ReadWriteLock` for Upgradable Read/Write, `Semaphores` can also be used at OS level, `Mutex`. Collections are **not** Thread-safe.

5. THREAD POOLS

Threads do have a cost. Many threads **slow down** the system. There is also a **Memory Cost**, because there is a stack for each thread. **Recycle** threads for multiple tasks to avoid this.

Tasks: Define **potentially parallel** work packages. Passive objects describing the functionality.

Thread Pool: Tasks are queued. A much smaller number of **working threads** grab tasks from the queue and execute them. A task must run to completion before a thread can grab a new one.

Scalable Performance: Performance problems with tasks run faster on parallel machines. This allows the exploitation of parallelism without thread pools. The number of threads can be **adapted** to the system. (Rule of thumb: # of Worker Threads = # processors + 1 (pending I/O calls))

Any task must **complete execution** before its worker thread is free to grab another task. Exception: nested tasks.

Advantages: Limited number of threads (Too many threads slow down the system or exceed available memory), **Thread recycling** (avoid thread creation and release), **Higher level of abstraction** (Disconnect task description from execution), **Number of threads configurable** on a per-system basis.

Limitations: Task must not wait for each other (nested sub-tasks), results in deadlocks (if one task `Ti` is waiting for something the `Tj` behind him in the Queue should wait, but `Tj` waits for `Ti` to finish, a deadlock occurs)

5.1. JAVA THREAD POOL

```
Task Launch
var threadPool = new ForkJoinPool(N);
Future<Integer> future = threadPool.submit(() -> { // submit task into pool
    int value = ...; // long calculation // return value; });
int value = ...; // long calculation // return value; }
```

5.1.1. Future<T>

Represents a **Future result** that is to be computed (asynchronous). Acts as proxy for the result that may be not available yet because the task has not finished. Usage: `var submit()` (submit task into pool), `launches task`, `get()` (waits if necessary for computation to complete and then returns its result) and `.cancel()` (Attempts to cancel execution of this task, removes it from queue). Task ends when an unhandled exception occurs. It is included in the `ExecutionException` thrown by `get()`.

5.1.2. Fire and Forget

Task is started **without retrieving results** later (submit() without get()). Task is run, but Exceptions will not get caught.

5.1.3. Count Prime Numbers

```
// Sequential
int counter = 0; for (int n = 2; n <= N; n++) { if (isPrime(n)) { counter++; }
// Parallel and Recursive
class CountPrimes extends RecursiveTask<Integer> { //RecursiveAction: void function
    private final int lower, upper;
    public CountPrimes(int lower, int upper) {
        { this.lower = lower; this.upper = upper; }
        protected Integer compute() {
            if (lower == upper) { return 0; }
            if (lower + 1 == upper) { return isPrime(lower) ? 1 : 0; }
            int middle = (lower + upper) / 2;
            var left = new CountTask(lower, middle);
            var right = new CountTask(middle, upper);
            left.fork(); right.fork();
            return right.join() + left.join();
        }
    }
}
```

`int result = new CountTask(2, N).invoke(); // invokeAll() to start multiple tasks`

5.1.4. PairwiseSum (recursive)

```
class PairwiseSum extends RecursiveAction {
    private final int[] array;
    private final int lower, upper;
    private static final int THRESHOLD = 1; // configurable
    public PairwiseSum(int[] array, int lower, int upper) {
        this.array = array; this.lower = lower; this.upper = upper;
    }
    protected void compute() {
        if (upper - lower > THRESHOLD) {
            int middle = (lower + upper) / 2;
            invokeAll(
                new PairwiseSum(array, lower, middle),
                new PairwiseSum(array, middle, upper));
        } else {
            for (int i = lower; i < upper; i++) {
                array[2*i] += array[2*i+1]; array[2*i+1] = 0;
            }
        }
    }
}
```

5.1.5. Work Stealing Thread Pool

Jobs get submitted into the **global queue**, which distributes the jobs to the **local queues** of each worker thread. If one thread has no work left, it can **"steal" work from another thread's** local queue instead of the global queue. This **distributes** the scheduling work over idle processors.

5.2. JAVA FORK JOIN POOL

Special Features: Fire-and-forget tasks may not finish, worker threads run as daemon threads. Automatic degree of parallelism (Default: As much worker threads as Processors).

5.3. .NET TASK PARALLEL LIBRARY (TPL)

Preferred way to write multi-threaded and parallel code. Provides public types and APIs in System.Threading and System.Threading.Tasks namespaces. **Efficiently duplicate thread pool** (tasks are queued to the ThreadPools, supports algorithms to provide load balancing, tasks are lightweight), **has multiple abstraction layers** (Task Parallelization: Use tasks explicitly, Data Parallelization: use parallel statements and queries using tasks implicitly), Asynchronous Programming and PLINQ.

```
// Task with return value in C#
Task<int> task = Task.Run(() -> {
    int total = ...; // some calculation
    return total;
});
Console.WriteLine(task.Result); // Blocks until task is done and returns the result
```

```
// Waited Task
var task = Task.Run(() => {
    var left = Task.Run(() => Count(leftPart));
    var right = Task.Run(() => Count(rightPart));
    return left.Result + right.Result;
});
static Task<int> Count(int... parts) { ... }
```

5.3.1. Parallel Statements in C#

Execute independent statements **potentially in parallel** (Ensure a task for each item, implicit barrier at the end).

```
Parallel.Invoke(
    () => MergeSort(l, m),
    () => MergeSort(m, r)
);
```

Execute **loop-bodies potentially in parallel** (Ensure a task for each item, implicit barrier at the end).

```
Parallel.ForEach(list,
    file => Convert(file));
Parallel.For(0, array.Length,
    i => DoCompute(array[i]));
```

Parallel Loop Partitioning: Loop with lots of quickly executing bodies. It would be inefficient to execute each iteration in parallel. Instead, TPL **automatically groups multiple bodies** into a single task. There are 4 kinds of **partitioning schemes**: Range (equally sized), Chunk (partitions start small and grow bigger), Stripe (n-sized partitions, where n = Small number, sometimes 17) and Hash partitioning (default: if input is indistinguishable then range partitioning, else chunk partitioning).

5.3.2. PLINQ

LINQ: Set of technologies based on the integration of SQL-like query capabilities directly into C#. LINQ is a parallel implementation of LINQ. Benefits from simplicity and readability of LINQ with the power of parallel programming by creating segments from its data. Analog to Java Stream API.

`from book in bookCollection.AsParallel() where book.Title.Contains("Concurrency") select book ISBN` // Random order

`first book in InputList.AsParallel().AsOrdered() select InputPrime(number)`

// Maintains order but is slower

5.3.3. Thread Injection

TPL adds new worker threads **at runtime** every time a work item completes or every 500ms.

Hit Climbing Algorithm: Maximize throughput while using as few threads as possible. Measures throughput & varies number of worker threads. Avoids deadlock with task-dependencies (but implicitly does not designed for this. Deadlocks with ThreadPools/WorkerThreads() are still possible). We should keep parallel tasks short to better profit from this automatic performance improvement.

6. ASYNCHRONOUS PROGRAMMING

Unnecessary Synchrony: Blocking method calls are often used without need (long running calculations on CDs, database or file accesses). With an **asynchronous call**, other work can continue while waiting on the result of the long operation.

`var task = Task.Run(LongOperation); /* other work // int result = task.Result;`

Kind of Asynchronism: *Call-centric* ("pass" calls waits for the task end and gets the result, blocking call), *Call-centric* ("pass" task hands over the result directly to successor / followup task).

Task Continuations: Define task object which is linked to the end of the predecessor task.

```
// C# .NET
Task CompleteFuture
.Run(task)
    .supplyAsync(() -> LongOp) // RunAsync for return void
    .ContinueWith(task2)
    .thenAsync(async v -> 2 * v) // returns value
    .ContinueWith(task3); // thenAcceptAsync(v -> ... .println(v)); // returns void
```

Multi-Continuation: Continue when all tasks are finished: `Task.WhenAll(task1, task2).ContinueWith(continuation);`

Continue when any one of the tasks are finished (other threads get lost after first thread is done): `Task.WhenAny(task1, task2).ContinueWith(continuation);`

(Exception in fire & forget task get ignored, i.e. `Task.Run(() => { ...; throw e; })`)

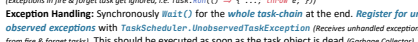
Exceptions Handling: Synchronously `Wait()` for the whole task-chain at the end. Register for unobserved exceptions with `TaskScheduler.UnobservedTaskException` (Receives unhandled exceptions from fire & forget tasks). This should be executed as soon as the task object is dead (Garbage Collector).

Java CompletableFuture: Modern asynchronous programming in Java. Also has **Multi-Continuation** with `CompletableFuture.allOf(Future1, future2) and CompletableFuture.anyOf(...)`

6.1. NON-BLOCKING GUI'S

Use case: If a UI is doing a long task, it should not freeze.

GUI Thread Model: Only single-threaded (Only a special UI-thread is allowed to access UI-components). The UI thread loops to process the event queue.



GUI Premise: No long operations in event, or else blocks UI. No access to UI-elements by other threads, or else incorrect (Exception in A.NET & Android, Race Condition in Java Swing).

6.1.1. Non-Blocking UI Implementation

```
// C# .NET
void buttonClick(...) {
    var url = textBox.Text;
    Task.Run(() => {
        var text = download(url); // to worker thread
        Dispatcher.InvokeAsync(
            () => {
                Label.Context = text;
            });
    });
}
```

Java InvokeLater: To be executed **asynchronously** on the event dispatching thread. Should be used when an **application thread** needs to **update the GUI**

